



Ramses Team

Members

Romain Teyssier (PU astro & PACM)

Jonathan Palafoutas (PU PACM)

Robel Geda (PU astro)

James Sunseri (PU astro)

Rayan Ait Ekioui (PU astro)


Takanori Akieda-Codron (PU astro)

Mentors

Matt Coleman

(PU Research Computing)

Brent Leback (Nvidia)






RAMSES (mini-ramses)

Adaptive Mesh Refinement for radiative MHD with self-gravity and PIC for astrophysical fluid flows

- PDE solvers using cell-by-cell adaptive grids
- MPI
- Modern Fortran
- Godunov Finite Volume Scheme
- Nvidia Fortran (CUDA)

<https://bitbucket.org/rteyssie/mini-ramses/src/develop/>




Goals

- First version written by Bob Caddy for a unigrid on a single GPU
- Extend to unigrid with multiple GPUs
- Profile the resulting code with Nsight Systems and Nsight Compute
- Optimize data transfer exploiting concurrency
- Explore next steps: AMR with multiple GPUs


Comparing kernel size, compute time, and memory

$$\frac{\# \text{ updates}}{\# \text{ kernel elements}} = \frac{2 \times 2 \times 2}{4 \times 4 \times 4} = \frac{1}{8}$$


$$= \frac{1}{3 \times 3 \times 3} = \frac{1}{27}$$




(ncentral=2)



(ncentral=1)




(ncentral=2)

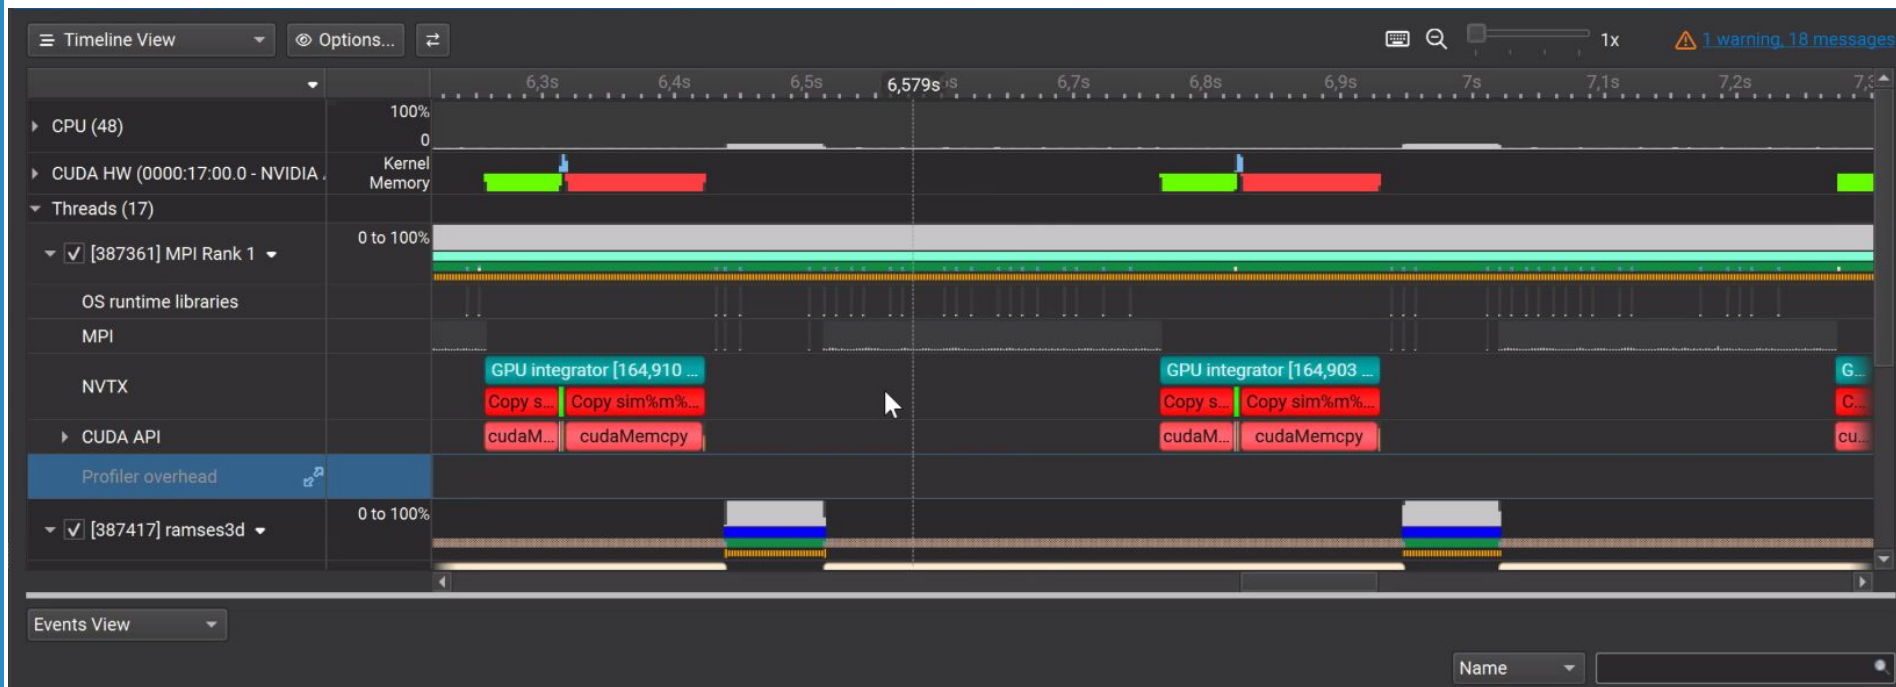


(ncentral=1)

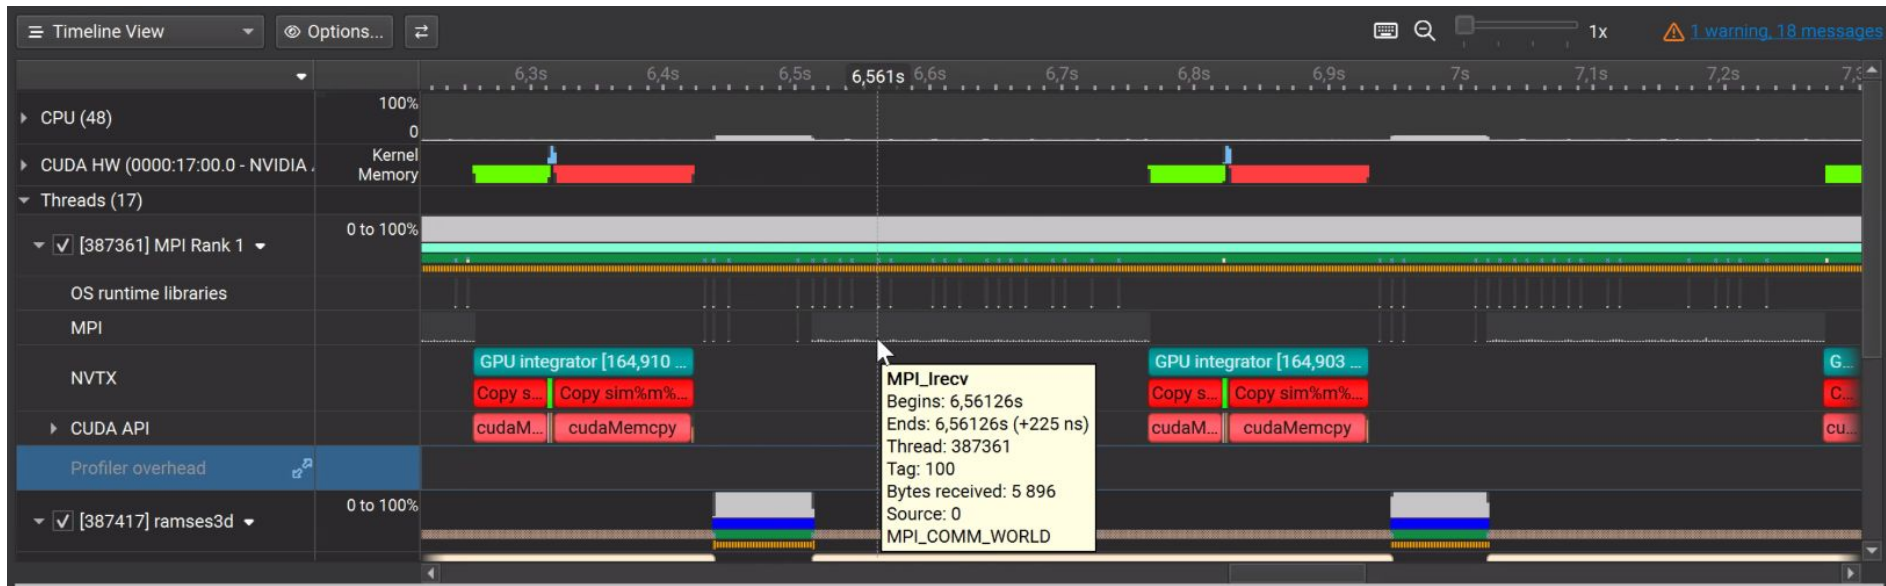
$$\frac{\text{old exec time}}{\text{new exec time}} = 3.2 \approx 27/8$$

Memory

- Shared memory per block for ncentral=**1**: **12 KB**
- Shared memory per block for ncentral=**2**: **40 KB**
- Shared memory per block for ncentral=4: **200 KB** 



GPU integrator is 4x faster due to proper parallelization.



GPU integrator is 4x faster due to proper parallelization.